



BAFAR

Infrastructure Design

Developed by Rayanesh Mobin Hooshyar



Esfand 1403

1. Floating IP Implementation

➤ Overview:

- The floating IP is implemented using **Keepalived** software and **VRRP (Virtual Router Redundancy Protocol)** . This setup ensures high availability by directing all incoming traffic to the floating IP address, which acts as a single entry point for clients.

➤ Failover Mechanism:

- In the event of a failure or outage on the **master server** , responsibilities are automatically reassigned to an alternative server with a lower priority designation.
- Keepalived continuously monitors the health of the master server and triggers failover if necessary, ensuring minimal downtime.

2. Load Balancing System

➤ Architecture:

The load balancing system is positioned behind the floating IP and operates as follows:

1. **Floating IP Layer** : All incoming traffic is directed to the floating IP.
2. **Swarm Networking Layer** : Traffic is managed by Docker Swarm's built-in networking, ensuring seamless communication between containers and services.
3. **Nginx Reverse Proxy Layer** : Within the Swarm cluster, an **Nginx reverse proxy** forwards traffic to the targeted service based on predefined rules.
4. **Targeted Service Layer** : Nginx directs traffic to the appropriate backend services running inside the Swarm cluster.

➤ Key Features:

- **High Availability** : Automatic failover ensures continuous service availability.
- **Scalability** : Dynamic scaling capabilities adapt to changing traffic demands.
- **Performance Optimization** : Efficient traffic management reduces latency and improves response times

3. MariaDB Cluster

➤ Databases: Replication Design

- **Replication Type** : Master/Slave with GTID-based asynchronous replication.
- **Binlog Format** : Row-based binary logging ensures accurate data replication.

➤ MaxScale Load Balancer

- **Overview** : All database traffic is proxied through **MaxScale** , ensuring optimal distribution of database requests.
- **Location** : MaxScale resides within the Swarm cluster, maintaining high availability and accessibility from the internal network.

3. MariaDB Cluster

➤ Read/Write Split

- **Write Queries** : Directed to the master server for data integrity and consistency.
- **Read Queries** : Dispatched to synchronized slave servers for load balancing and reduced latency.

➤ Failover Mechanism

- If the master server becomes unresponsive for 90 seconds, MaxScale automatically promotes one of the slave servers to the master role.
- Remaining slave servers are reconfigured to synchronize with the new master.
- If the original master recovers, it is reintegrated into the system as a slave.

4. Web Services

➤ Overview

Web services and their dependencies are hosted within the Docker Swarm cluster. Dependencies include:

- **MariaDB** : Interfaced with MaxScale for database access.
- **Redis** : Used for caching purposes.
- **Shared Storage** : Ensures accessible and consistent data storage across services

➤ Shared Storage Configuration

- **GlusterFS** : Configured to manage persistent data effectively.
- The /mnt path is shared among all servers in the Swarm cluster.

➤ Failover Strategy

- In the event of a service outage on a server, affected services are automatically rescheduled and redistributed to other operational servers within the cluster.

➤ Scalability

- Swarm's inherent scalability allows container services to scale dynamically based on demand.

5. Deployment Strategy

➤ CI/CD Implementation

- The Continuous Integration/Continuous Deployment (CI/CD) strategy leverages **Gitlab-CI** and **Ansible** for automated deployments.

➤ Gitlab-CI Workflow

1. **Initial Stages** : Executes 'Test' and 'Build' processes.
2. **Deployment Trigger** : Activates a deployment webhook exclusively for tag events prefixed with production .

➤ Ansible Workflow

- Ansible invokes deployment processes via the webhook triggered by Gitlab-CI.
- An **Nginx server** is configured to return deployment status, ensuring transparency and traceability during the deployment phase

6. Monitoring and Alerting

➤ Overview

- To ensure the health and performance of the infrastructure, we use a robust monitoring stack consisting of **Prometheus**, **Grafana**, and **Alertmanager**. This setup provides real-time insights into system metrics, enables proactive issue detection, and facilitates timely alerts for critical situations.

➤ Monitoring Components

1. Prometheus

- **Role:** Prometheus serves as the primary time-series database and monitoring engine.
- **Metrics Collected:**
 - **Incoming HTTP Traffic:** Monitors request rates, response times, and error rates for all services.
 - **Host Resources:** Tracks CPU usage, memory consumption, disk I/O, and network bandwidth across all servers.
 - **Container Metrics:** Monitors container-level resource utilization (CPU, memory, disk, and network) via Docker Swarm integration.

6. Monitoring and Alerting

2. Grafana

- **Role:** Grafana is used to visualize Prometheus metrics through customizable dashboards.
- **Key Dashboards:**
 - **HTTP Traffic Dashboard** : Displays trends in incoming requests, response times, and error rates.
 - **Host Resource Dashboard** : Provides an overview of server health, including CPU, memory, disk, and network utilization.
 - **Container Health Dashboard** : Shows the status and resource usage of all containers in the Swarm cluster.\

6. Monitoring and Alerting

3. Alertmanager

- **Role:** Alertmanager handles alerts sent by Prometheus and routes them to appropriate notification channels.
- **Alerting Rules :**
 - **Critical Situations :**
 - High CPU or memory usage on any host exceeding predefined thresholds.
 - Disk space utilization above 85%.
 - Network outages or high latency detected in floating IP failover mechanisms.
 - Redis evictions due to insufficient memory.
 - **Service Degradation :**
 - Increased error rates in HTTP responses (e.g., 5xx errors).
 - Slow response times for critical services.
 - **Failover Events :**
 - Automatic notifications when a failover occurs in the floating IP or database cluster.

